

Visualization of the Data

After cleaning the dataset, I imported it and used it to visualize it to produce insights that will be of use and help to the client, Microsoft.

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```

```
In [2]: df = pd.read_csv('cleaned_data.csv',)
df
```

| | 0 | production_budget | worldwide_gross | original_title | start_year | runtime_minutes | genres |
|------|------|-------------------|-----------------|------------------------------|------------|-----------------|-------------------------|
| 0 | 0 | 80000000.0 | 247023808.0 | The Courier | 2012 | 95.0 | Action, Crime, Drama |
| 1 | 1 | 80000000.0 | 226739416.0 | Pelé: Birth of a Legend | 2016 | 107.0 | Biography, Drama, Sport |
| 2 | 2 | 80000000.0 | 168311558.0 | Bibliothèque Pascal | 2010 | 105.0 | Drama |
| 3 | 3 | 80000000.0 | 241200000.0 | The Experiment | 2010 | 96.0 | Drama, Thriller |
| 4 | 4 | 80000000.0 | 221468935.0 | Stadilaista tangoa etsimässä | 2010 | 90.0 | Documentary, Music |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2110 | 2172 | 7000.0 | 841926.0 | Pet | 2016 | 94.0 | Horror, Thriller |
| 2111 | 2173 | 7000.0 | 71644.0 | Stag Night of the Dead | 2010 | 81.0 | Action, Comedy, Horror |

```
In [3]: df = df.drop(columns=['Unnamed: 0'])
```

```
In [4]: df.columns
```

```
Out[4]: Index(['production_budget', 'worldwide_gross', 'original_title', 'start_year',
              'runtime_minutes', 'genres'],
              dtype='object')
```

```
In [5]: # Calculate the correlation matrix
correlation = df['production_budget'].corr(df['worldwide_gross'])

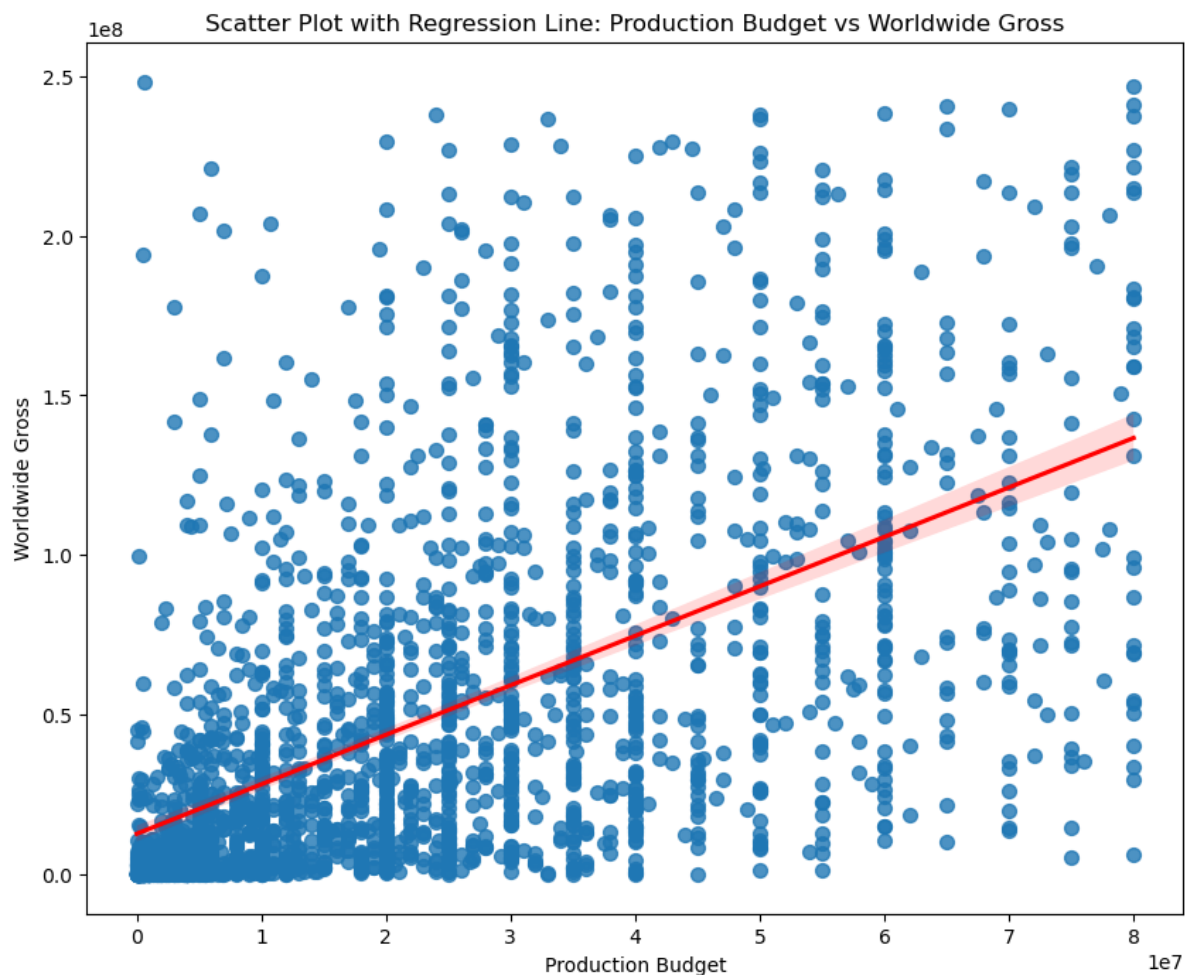
# Display the correlation coefficient
print(f"Correlation between 'production_budget' and 'worldwide_gross': {correlation:.2f}")
```

Correlation between 'production_budget' and 'worldwide_gross': 0.57

```
In [6]: # Exploring the relationship between Production Budget and Worldwide Gross using a scatter plot and a trend line
plt.figure(figsize=(10, 8))
sns.regplot(x='production_budget', y='worldwide_gross', data=df, scatter_kws={'s': 50}, line_kws={'color': 'red'})

# Set Labels and title
plt.title('Scatter Plot with Regression Line: Production Budget vs Worldwide Gross')
plt.xlabel('Production Budget')
plt.ylabel('Worldwide Gross')

# Display the plot
plt.show();
# As Production budget increases, so does the Worldwide Gross increase
```

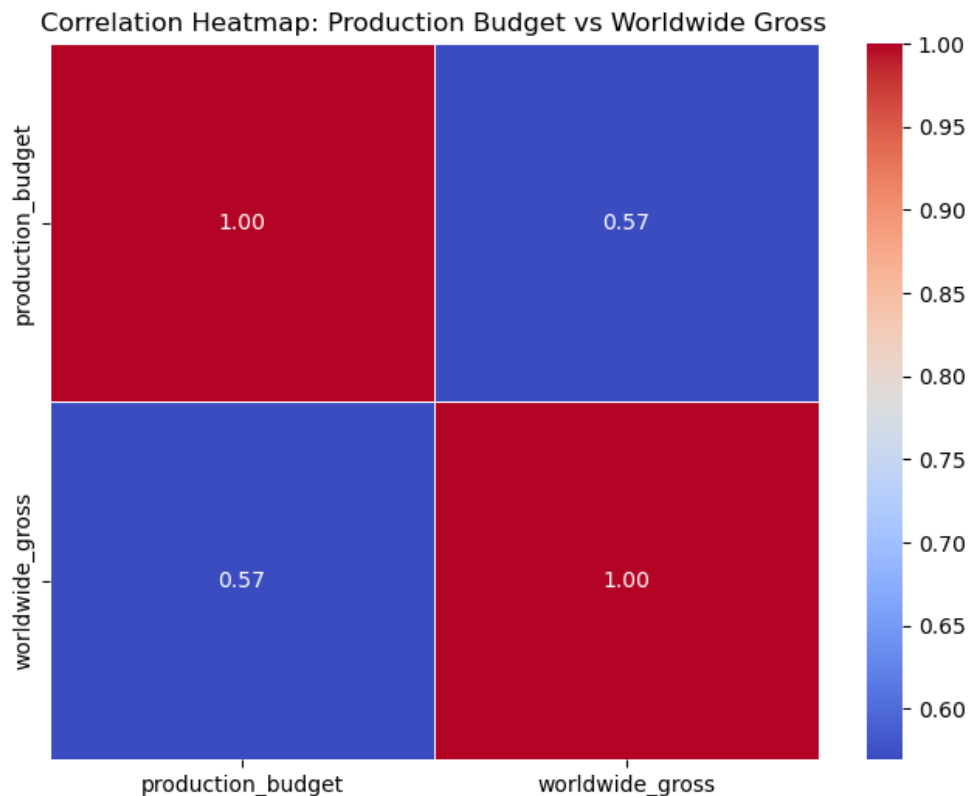


```
In [7]: # Creating a DataFrame with only production_budget and worldwide_gross to understand their correlation
selected_columns = ['production_budget', 'worldwide_gross']
subset_df = df[selected_columns]

# Calculate the correlation matrix for the selected columns
correlation_matrix_subset = subset_df.corr()

# Create a heatmap using Seaborn
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix_subset, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)

# Set labels and title
plt.title('Correlation Heatmap: Production Budget vs Worldwide Gross')
# Display the plot
plt.show()
# A positive correlation exists between Production Budget and Worldwide Gross, which is a good thing
# Production budget Leads to an increase in Worldwide Gross
```



```
In [8]: print(df.columns)

Index(['production_budget', 'worldwide_gross', 'original_title', 'start_year',
      'runtime_minutes', 'genres'],
      dtype='object')
```

```
In [9]: # Determining the top 3 genres with high worldwide gross generation
# Split genres into separate rows
df_genres = df['genres'].str.split(',', expand=True).stack().reset_index(level=1, drop=True).to_frame()
df_expanded = df_genres.join(df.drop('genres', axis=1))


# Group by genre and calculate the sum of worldwide_gross for each genre
genre_worldwide_gross = df_expanded.groupby('genre')['worldwide_gross'].sum()

# Find the top 3 genres
top_genres = genre_worldwide_gross.nlargest(3)

print("Top 3 genres with high worldwide gross:")
print(top_genres)
# Top 3 genres with high worldwide gross generation are Documentary, Drama and Comedy.
```

```
Top 3 genres with high worldwide gross:
genre
Documentary    4.424609e+10
Drama          3.062846e+10
Comedy         1.343596e+10
Name: worldwide_gross, dtype: float64
```

```
In [10]: # Create a bar plot to measure genres with highest worldwide gross.
plt.figure(figsize=(10, 6))
genre_worldwide_gross.sort_values(ascending=False).plot(kind='bar', color='skyblue')
plt.title('Worldwide Gross by Genre')
plt.xlabel('Genre')
plt.ylabel('Worldwide Gross')
plt.show();
# According to genre classification by Worldwide Gross, Documentaries seem to be pooling massive Worl
# by drama and comedy
```



```
In [11]: # Calculate the total number of movies per genre
genre_counts = df['genres'].value_counts().reset_index()
genre_counts.columns = ['Genre', 'Total Movies']

# Rank genres based on the total number of movies
genre_counts['Rank'] = genre_counts['Total Movies'].rank(ascending=False)
genre_counts.head(6)
```

Out[11]:

| | Genre | Total Movies | Rank |
|---|--------------|--------------|------|
| 0 | Documentary | 717 | 1.0 |
| 1 | Drama | 321 | 2.0 |
| 2 | Comedy | 116 | 3.0 |
| 3 | Comedy,Drama | 62 | 4.0 |
| 4 | Horror | 60 | 5.0 |
| 5 | Thriller | 49 | 6.0 |


```
In [12]: # Calculate the correlation matrix
correlation_matrix = df[['runtime_minutes', 'production_budget', 'worldwide_gross']].corr()

# Display the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)
```

```
Correlation Matrix:
runtime_minutes  runtime_minutes  production_budget  worldwide_gross
runtime_minutes      1.000000      -0.022727      -0.002747
production_budget  -0.022727      1.000000      0.569324
worldwide_gross    -0.002747      0.569324      1.000000
```

```
In [13]: # Determining the correlation between runtime and Production Budget and Worldwide Gross
# Create a heatmap using Seaborn
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)

# Set Labels and title
plt.title('Correlation Heatmap: Runtime, Production Budget, Worldwide Gross')
plt.show()
# There is a negative correlation between runtime minutes and Worldwide Gross, and similarly, a negat
# exists between runtime minutes and Production Budget.
```



```
In [14]: # Calculating the profit margins based on worldwide gross and production budget and add it to the Dat
df['profit_margin'] = ((df['worldwide_gross'] - df['production_budget']) / df['worldwide_gross']) * 1
df
```

Out[14]:


| | production_budget | worldwide_gross | original_title | start_year | runtime_minutes | genres | profit_margi |
|------|-------------------|-----------------|------------------------------|------------|-----------------|-----------------------------|--------------|
| 0 | 80000000.0 | 247023808.0 | The Courier | 2012 | 95.0 | Action,Crime,Drama | 67.61445 |
| 1 | 80000000.0 | 226739416.0 | Pelé: Birth of a Legend | 2016 | 107.0 | Biography,Drama,Sport | 64.71720 |
| 2 | 80000000.0 | 168311558.0 | Bibliothèque Pascal | 2010 | 105.0 | Drama | 52.46909 |
| 3 | 80000000.0 | 241200000.0 | The Experiment | 2010 | 96.0 | Drama,Thriller | 66.83250 |
| 4 | 80000000.0 | 221468935.0 | Stadilaista tangoa etsimässä | 2010 | 90.0 | Documentary,Music | 63.87755 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2110 | 7000.0 | 841926.0 | Pet | 2016 | 94.0 | Horror,Thriller | 99.16857 |
| 2111 | 7000.0 | 71644.0 | Stag Night of the Dead | 2010 | 81.0 | Action,Comedy,Horror | 90.22946 |
| 2112 | 7000.0 | 900.0 | Mr. Nice | 2010 | 121.0 | Biography,Comedy,Crime | -677.77777 |
| 2113 | 6000.0 | 240495.0 | Marley | 2012 | 144.0 | Biography,Documentary,Music | 97.50514 |
| 2114 | 5000.0 | 1338.0 | Welcome to the Rileys | 2010 | 110.0 | Drama | -273.69207 |

2115 rows × 7 columns

```
In [15]: # Categorizing movies as profitable or non-profitable
df['profitable'] = df['profit_margin'] > 0

# Count the number of profitable and non-profitable movies
profitable_counts = df['profitable'].value_counts()

# Create a bar plot
plt.figure(figsize=(8, 6))
profitable_counts.plot(kind='bar', color=['green', 'red'])
plt.title('Number of Profitable vs Non-Profitable Movies')
plt.xlabel('Profitability')
plt.ylabel('Number of Movies')
plt.xticks([0, 1], ['Profitable', 'Non-Profitable'], rotation=0)
plt.show()
```



In [16]:

Determining the ROI: Calculating the ROI using production_budget and worldwide_gross then adding it to df
df['roi'] = ((df['worldwide_gross'] - df['production_budget']) / df['production_budget']) * 100
df

Out[16]:

| | production_budget | worldwide_gross | original_title | start_year | runtime_minutes | genres | profit_margin |
|------|-------------------|-----------------|------------------------------|------------|-----------------|-----------------------------|---------------|
| 0 | 80000000.0 | 247023808.0 | The Courier | 2012 | 95.0 | Action,Crime,Drama | 67.61445 |
| 1 | 80000000.0 | 226739416.0 | Pelé: Birth of a Legend | 2016 | 107.0 | Biography,Drama,Sport | 64.71720 |
| 2 | 80000000.0 | 168311558.0 | Bibliothèque Pascal | 2010 | 105.0 | Drama | 52.46909 |
| 3 | 80000000.0 | 241200000.0 | The Experiment | 2010 | 96.0 | Drama,Thriller | 66.83250 |
| 4 | 80000000.0 | 221468935.0 | Stadilaista tangoa etsimässä | 2010 | 90.0 | Documentary,Music | 63.87755 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2110 | 7000.0 | 841926.0 | Pet | 2016 | 94.0 | Horror,Thriller | 99.16857 |
| 2111 | 7000.0 | 71644.0 | Stag Night of the Dead | 2010 | 81.0 | Action,Comedy,Horror | 90.22946 |
| 2112 | 7000.0 | 900.0 | Mr. Nice | 2010 | 121.0 | Biography,Comedy,Crime | -677.77777 |
| 2113 | 6000.0 | 240495.0 | Marley | 2012 | 144.0 | Biography,Documentary,Music | 97.50514 |
| 2114 | 5000.0 | 1338.0 | Welcome to the Rileys | 2010 | 110.0 | Drama | -273.69207 |

2115 rows × 9 columns




```
In [17]: # Determining the return on investment by genre type to find out which genre should be considered for
# Splitting genres into separate rows
df_genres = df['genres'].str.split(',', expand=True).stack().reset_index(level=1, drop=True).to_frame
df_expanded = df_genres.join(df.drop('genres', axis=1))

# Create a bar plot with sorted ROI values (largest to smallest)
plt.figure(figsize=(10, 6))
sns.barplot(x='genre', y='roi', data=df_expanded, ci=None, palette='viridis', order=df_expanded.groupby('genre')['roi'].median().sort_values(ascending=False).index[::-1])
plt.title('ROI by Genre')
plt.xlabel('Genre')
plt.ylabel('Return on Investment (%)')
plt.xticks(rotation=45, ha='right')
plt.show()
# Top 3 genres that MS can consider to invest in are Sci-Fi, Crime, and Mystery, where Mystery movies
```


C:\Users\DELL\AppData\Local\Temp\ipykernel_20568\3891381107.py:8: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x='genre', y='roi', data=df_expanded, ci=None, palette='viridis', order=df_expanded.groupby('genre')['roi'].median().sort_values(ascending=False).index[::-1])
```




```
In [22]: # Determining the relationship between movie runtime minutes and ROI by scatter plot analysis
plt.figure(figsize=(10, 6))
plt.scatter(df['runtime_minutes'], df['roi'], color='purple')
plt.title('Runtime Minutes vs Return on Investment (ROI)')
plt.xlabel('Runtime Minutes')
plt.ylabel('Return on Investment')
plt.show()
# There is a higher ROI on movies whose runtime minutes are within 80 - 120 minutes
```

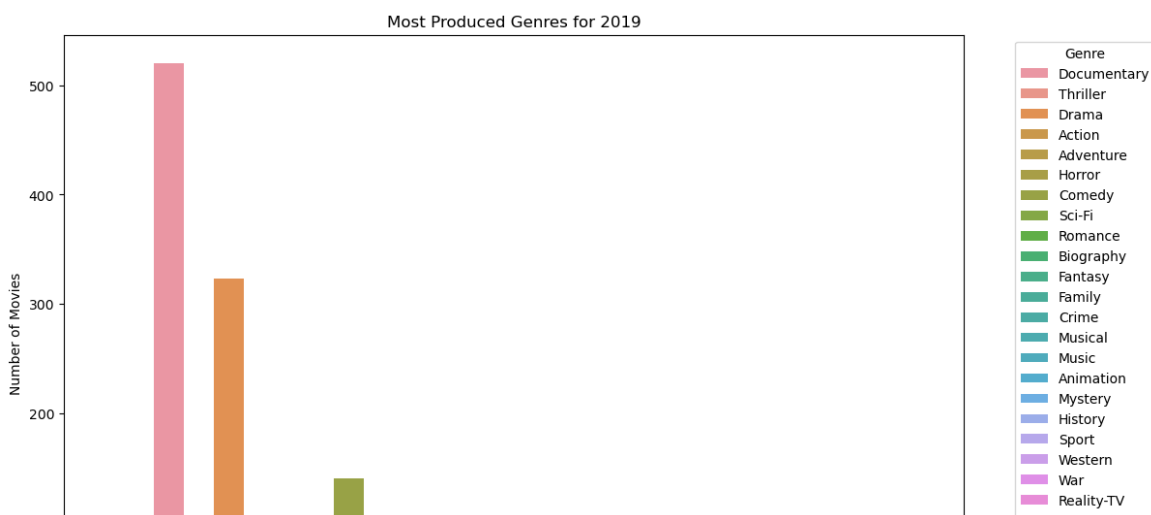


```
In [19]: # Determining which year had the highest production of genres
# Split genres into separate rows
df_genres = df['genres'].str.split(',', expand=True).stack().reset_index(level=1, drop=True).to_frame()
df_expanded = df_genres.join(df.drop('genres', axis=1))

# Choose the topmost year
top_year = df['start_year'].value_counts().idxmax()

# Filter data for the top year
df_top_year = df_expanded[df_expanded['start_year'] == top_year]

# Create a countplot
plt.figure(figsize=(12, 8))
sns.countplot(data=df_top_year, x='start_year', hue='genre')
plt.title(f'Most Produced Genres for {top_year}')
plt.xlabel('Year')
plt.ylabel('Number of Movies')
plt.legend(title='Genre', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```



```
In [20]: # Determining the relationship between movie runtime minutes and ROI by scatter plot analysis  
plt.figure(figsize=(10, 6))  
plt.scatter(df['runtime_minutes'], df['production_budget'], color='blue')  
plt.title('Runtime Minutes vs Production Budget')  
plt.xlabel('Runtime Minutes')  
plt.ylabel('Production Budget')  
plt.show();  
# Shorter films have a low production budget  
# Longer films have a high production budget  
# Most of the films lying between 80 - 120 mins have an relatively average production budget
```




```
In [21]: # Determining the trend in profit margin of the movies across the years
# Create a line plot to show the profit margin trend by year
plt.figure(figsize=(12, 8))
sns.lineplot(data=df, x='start_year', y='profit_margin', marker='o', ci=None)
plt.title('Profit Margin Trend by Year')
plt.xlabel('Year')
plt.ylabel('Profit Margin (%)')
plt.show()
# A stability on profit margin is seen between years 2010 to 2018. A sharp drop in profit margin note
# and then picks at 2022. By assumption, the Covid-19 pandemic may have been the major cause of the s
# However, MS can proceed to invest in movie production as the trend shows positive increment accordi
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_20568\466270572.py:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.lineplot(data=df, x='start_year', y='profit_margin', marker='o', ci=None)
```



These are the usable visualizations of the data